

Linear Methods for Classification

Flavien De Bock

2026-08-26

```
rm(list = ls())
library(readr) ; library(tidyverse) ; library(corrplot) ; library(MASS)


train <- read.csv("vowel_train.csv", sep = ",")
train$x.10. <- gsub("\\\\", "", train$x.10.)
train$x.10. <- as.numeric(gsub("}", "", train$x.10.))
train$x.10 <- train$x.10. ; train$x.10. <- NULL

test <- read.csv("vowel_test.csv", sep = ",")
test$x.10. <- gsub("\\\\", "", test$x.10.)
test$x.10. <- as.numeric(gsub("}", "", test$x.10.))
test$x.10 <- test$x.10. ; test$x.10. <- NULL
head(train) ; summary(train)
```

```
##   y    x.1    x.2    x.3    x.4    x.5    x.6    x.7    x.8    x.9    x.10
## 1 1 -3.639 0.418 -0.670 1.779 -0.168 1.627 -0.388 0.529 -0.874 -0.814
## 2 2 -3.327 0.496 -0.694 1.365 -0.265 1.933 -0.363 0.510 -0.621 -0.488
## 3 3 -2.120 0.894 -1.576 0.147 -0.707 1.559 -0.579 0.676 -0.809 -0.049
## 4 4 -2.287 1.809 -1.498 1.012 -1.053 1.060 -0.567 0.235 -0.091 -0.795
## 5 5 -2.598 1.938 -0.846 1.062 -1.633 0.764 0.394 -0.150 0.277 -0.396
## 6 6 -2.852 1.914 -0.755 0.825 -1.588 0.855 0.217 -0.246 0.238 -0.365

##           y           x.1           x.2           x.3
## Min.      : 1    Min.      :-5.211    Min.      :-1.2740   Min.      :-2.48700
## 1st Qu.: 3    1st Qu.: -3.923    1st Qu.: 0.9167    1st Qu.: -0.94550
## Median : 6    Median : -3.097    Median : 1.7330    Median : -0.50250
## Mean      : 6    Mean      : -3.167    Mean      : 1.7353    Mean      : -0.44800
## 3rd Qu.: 9    3rd Qu.: -2.512    3rd Qu.: 2.4038    3rd Qu.: 0.04925
## Max.      :11    Max.      : -0.941    Max.      : 5.0740    Max.      : 1.41300
##           x.4           x.5           x.6           x.7
## Min.      : -1.4090    Min.      : -2.1270    Min.      : -0.8360    Min.      : -1.53700
## 1st Qu.: -0.0835    1st Qu.: -0.9307    1st Qu.: 0.1085    1st Qu.: -0.29700
## Median : 0.4565    Median : -0.4170    Median : 0.5275    Median : 0.04000
## Mean      : 0.5250    Mean      : -0.3893    Mean      : 0.5850    Mean      : 0.01748
## 3rd Qu.: 1.1640    3rd Qu.: 0.1155    3rd Qu.: 1.0097    3rd Qu.: 0.34800
## Max.      : 2.1910    Max.      : 1.8310    Max.      : 2.3270    Max.      : 1.40300
##           x.8           x.9           x.10
## Min.      : -1.29300    Min.      : -1.6130    Min.      : -1.68000
## 1st Qu.: -0.01825    1st Qu.: -0.6737    1st Qu.: -0.50700
## Median : 0.47700    Median : -0.2550    Median : -0.08250
## Mean      : 0.41739    Mean      : -0.2681    Mean      : -0.08457
## 3rd Qu.: 0.86125    3rd Qu.: 0.1375    3rd Qu.: 0.30100
## Max.      : 1.67300    Max.      : 1.3090    Max.      : 1.39600
```

```
corr_matrix <- cor(train[,-1]) # correlation matrix
corrplot(corr_matrix, method = "square", addCoef.col = 'black') # associated plot
```



```
train[paste0("y_", 1:11)] <- lapply(1:11, function(k) { as.integer(train$y == k)})
test[paste0("y_", 1:11)] <- lapply(1:11, function(k) { as.integer(test$y == k)})

x_vars <- paste0("x.", 1:10) # string vector of the name of the features

# Regression Indicator Matrix -----

models <- vector("list", 11) # create an empty list to store the 11 regressions

for (k in 1:11) { # loop through the 11 groups
  y_var <- paste0("y_", k) # string vector of the name of group k

  formula_k <- as.formula(
    paste(y_var, "~", paste(x_vars, collapse = " + ")) # formula regress y_k on the features
  )

  models[[k]] <- lm(formula_k, data = train) # fit the reg
}

fitted_train <- as.data.frame(
  lapply(models, fitted) # retrieve the fitted values for each model
) # we store them in a data.frame
```

```
names(fitted_train) <- paste0("y_", 1:11) # rename the columns for clarity

# the predicted outcome is the one with the highest fitted values
prediction_train <- max.col(fitted_train, ties.method = "first")

# Training error rate is ~0.48
train_error_rate <- mean(train$y != prediction_train) ; train_error_rate
```

```
## [1] 0.4772727
```

```
# we obtain the fitted values of the model on the test data
fitted_test <- as.data.frame(
  lapply(models, function(model) predict(model, newdata = test))
)
```

```
names(fitted_test) <- paste0("y_", 1:11)
```

```
prediction_test <- max.col(fitted_test, ties.method = "first")
```

```
# the test error rate is ~0.67
```

```
test_error_rate <- mean(test$y != prediction_test) ; test_error_rate
```

```
## [1] 0.6666667
```

```
# LDA -----
```

```
lda_model <- lda(y ~ x.1 + x.2 + x.3 + x.4 + x.5 + x.6 + x.7 + x.8 + x.9 + x.10,
  data = train) ; lda_model
```

```
## Call:
```

```
## lda(y ~ x.1 + x.2 + x.3 + x.4 + x.5 + x.6 + x.7 + x.8 + x.9 +
##      x.10, data = train)
```

```
##
```

```
## Prior probabilities of groups:
```

```
##      1      2      3      4      5      6      7
## 0.09090909 0.09090909 0.09090909 0.09090909 0.09090909 0.09090909 0.09090909
##      8      9     10     11
## 0.09090909 0.09090909 0.09090909 0.09090909
```

```
##
```

```
## Group means:
```

```
##      x.1      x.2      x.3      x.4      x.5      x.6      x.7
## 1 -3.359563 0.0629375 -0.2940625 1.2033333 0.38747917 1.2218958 0.09637500
## 2 -2.708875 0.4906042 -0.5802292 0.8135000 0.20193750 1.0634792 -0.19091667
## 3 -2.440250 0.7748750 -0.7983958 0.8086667 0.04245833 0.5692500 -0.28006250
## 4 -2.226604 1.5258333 -0.8744375 0.4221458 -0.37131250 0.2483542 -0.01895833
## 5 -2.756312 2.2759583 -0.4657292 0.2253125 -1.03679167 0.3897917 0.23641667
## 6 -2.673542 1.7587708 -0.4745625 0.3505625 -0.66585417 0.4170000 0.16233333
## 7 -3.243729 2.4683542 -0.1050625 0.3964583 -0.98029167 0.1623125 0.01958333
## 8 -4.051333 3.2339792 -0.1739792 0.3965833 -1.04602083 0.1951875 0.08666667
## 9 -3.876896 2.3450208 -0.3668333 0.3170417 -0.39450000 0.8033750 0.02504167
## 10 -4.506146 2.6885625 -0.2849167 0.4695625 -0.03879167 0.6388750 0.13916667
## 11 -2.990396 1.4638750 -0.5098125 0.3716458 -0.38039583 0.7250417 -0.08339583
##      x.8      x.9      x.10
## 1 0.03710417 -0.62435417 -0.16162500
## 2 0.37381250 -0.51595833 0.08060417
```

```

## 3  0.20495833 -0.47827083  0.18187500
## 4  0.10714583 -0.32627083 -0.05375000
## 5  0.42462500 -0.20070833 -0.28070833
## 6  0.22925000 -0.20750000  0.05270833
## 7  0.76229167 -0.03027083 -0.12239583
## 8  0.82077083  0.10445833  0.02122917
## 9  0.73614583 -0.23183333 -0.14810417
## 10 0.38756250 -0.11102083 -0.27335417
## 11 0.50766667 -0.32750000 -0.22672917
##
## Coefficients of linear discriminants:
##          LD1          LD2          LD3          LD4          LD5          LD6
## x.1 -0.904263484 -1.0765784  0.2138645  0.45014309  0.01568403  1.3197512
## x.2  1.150256514 -0.3481586 -0.1219356  0.81974744 -0.07879069  1.2828586
## x.3  0.539113617  0.4775768  1.3152720 -0.18174531 -0.01428301  1.0421541
## x.4  0.024636593  0.6124198  0.3333835  0.89249728  1.19422174  0.2840921
## x.5 -0.007828209  1.6078817 -1.2374917  0.96925109  0.17984293  0.9791226
## x.6  0.708040285  1.4820258  0.8767861 -0.06187384 -0.84899844  1.9958303
## x.7  0.843500553  0.8675407  1.4607005  1.10382558  1.35657038  1.0703157
## x.8  1.305220749  0.9393954 -0.9967222 -1.05415855  0.85277476 -0.1212728
## x.9  0.965050787  0.5131795 -0.6927861  0.60778087 -0.04667900  0.5487891
## x.10 0.352677857  0.1517412 -0.7092612 -0.22606164  1.04293623  1.7009401
##          LD7          LD8          LD9          LD10
## x.1 -1.2563534  0.55370976 -0.21893931  0.01475366
## x.2 -0.8543263  0.85921988 -0.46773378 -0.20534101
## x.3 -1.3953387 -0.05664837 -0.58464728 -0.52218214
## x.4 -0.6079608  1.18671349 -0.26235352  0.22298959
## x.5 -1.5595892  0.12322124  0.38252584  0.05017199
## x.6 -0.2001485  0.54892997 -0.82835032  0.13158053
## x.7 -0.5029121  0.09927612  1.50030138  0.36516484
## x.8 -0.6617860  0.25729462  1.12742075  0.30135489
## x.9 -0.5244975 -1.05461865 -0.04187258  1.31351481
## x.10 0.5629667 -0.29898796  0.14972818 -0.23691329
##
## Proportion of trace:
##          LD1          LD2          LD3          LD4          LD5          LD6          LD7          LD8          LD9          LD10
## 0.5617 0.3518 0.0445 0.0191 0.0107 0.0083 0.0026 0.0011 0.0001 0.0001

lda_prediction_train <- predict(lda_model, newdata = train)
lda_prediction_test  <- predict(lda_model, newdata = test)

# training error rate is ~ 0.32
lda_train_error_rate <- mean(train$y != lda_prediction_train$class) ; lda_train_error_rate

## [1] 0.3162879

# test error rate is ~ 0.56
lda_test_error_rate <- mean(test$y != lda_prediction_test$class) ; lda_test_error_rate

## [1] 0.5562771

# QDA -----

qda_model <- qda(y ~ x.1 + x.2 + x.3 + x.4 + x.5 + x.6 + x.7 + x.8 + x.9 + x.10,
                 data = train) ; qda_model

```

```

## Call:
## qda(y ~ x.1 + x.2 + x.3 + x.4 + x.5 + x.6 + x.7 + x.8 + x.9 +
##      x.10, data = train)
##
## Prior probabilities of groups:
##      1      2      3      4      5      6      7
## 0.09090909 0.09090909 0.09090909 0.09090909 0.09090909 0.09090909 0.09090909
##      8      9     10     11
## 0.09090909 0.09090909 0.09090909 0.09090909
##
## Group means:
##      x.1      x.2      x.3      x.4      x.5      x.6      x.7
## 1  -3.359563 0.0629375 -0.2940625 1.2033333 0.38747917 1.2218958 0.09637500
## 2  -2.708875 0.4906042 -0.5802292 0.8135000 0.20193750 1.0634792 -0.19091667
## 3  -2.440250 0.7748750 -0.7983958 0.8086667 0.04245833 0.5692500 -0.28006250
## 4  -2.226604 1.5258333 -0.8744375 0.4221458 -0.37131250 0.2483542 -0.01895833
## 5  -2.756312 2.2759583 -0.4657292 0.2253125 -1.03679167 0.3897917 0.23641667
## 6  -2.673542 1.7587708 -0.4745625 0.3505625 -0.66585417 0.4170000 0.16233333
## 7  -3.243729 2.4683542 -0.1050625 0.3964583 -0.98029167 0.1623125 0.01958333
## 8  -4.051333 3.2339792 -0.1739792 0.3965833 -1.04602083 0.1951875 0.08666667
## 9  -3.876896 2.3450208 -0.3668333 0.3170417 -0.39450000 0.8033750 0.02504167
## 10 -4.506146 2.6885625 -0.2849167 0.4695625 -0.03879167 0.6388750 0.13916667
## 11 -2.990396 1.4638750 -0.5098125 0.3716458 -0.38039583 0.7250417 -0.08339583
##      x.8      x.9      x.10
## 1  0.03710417 -0.62435417 -0.16162500
## 2  0.37381250 -0.51595833 0.08060417
## 3  0.20495833 -0.47827083 0.18187500
## 4  0.10714583 -0.32627083 -0.05375000
## 5  0.42462500 -0.20070833 -0.28070833
## 6  0.22925000 -0.20750000 0.05270833
## 7  0.76229167 -0.03027083 -0.12239583
## 8  0.82077083 0.10445833 0.02122917
## 9  0.73614583 -0.23183333 -0.14810417
## 10 0.38756250 -0.11102083 -0.27335417
## 11 0.50766667 -0.32750000 -0.22672917

qda_prediction_train <- predict(qda_model, newdata = train)
qda_prediction_test  <- predict(qda_model, newdata = test)

# training error rate is ~ 0.01
qda_train_error_rate <- mean(train$y != qda_prediction_train$class) ; qda_train_error_rate

## [1] 0.01136364

# test error rate is ~ 0.53
qda_test_error_rate <- mean(test$y != qda_prediction_test$class) ; qda_test_error_rate

## [1] 0.5281385

```